

# AI Agent Pilot Scorecard

A scoring rubric for pilots where leaders need to compare agents on usefulness, reliability, cost, latency, and auditability.

## AUDIENCE

**AI platform teams and procurement reviewers**

## USE IN

**15 min**

## BEST NEXT STEP

**Turn answers into a challenge pack and gate.**

---

## Score Dimensions

- 1 Correct final outcome and task completion.
- 2 Tool selection quality and recovery after failed calls.
- 3 Cost per successful task, not raw token spend.

---

## Reliability Signals

- 1 Repeated runs on the same pack show stable behavior.
- 2 Failures cluster into named categories with owners.
- 3 Candidate keeps quality when inputs include messy real-world details.

---

## Evidence Quality

- 1 Replay explains the path from prompt to final answer.
- 2 Artifacts are downloadable and tied to the scorecard.
- 3 Security can review tool paths without reading raw logs for hours.

---

## Pilot Decision

- 1 Ship, extend pilot, or reject with the exact failed dimension.
- 2 Name the next workload to freeze if the pilot advances.
- 3 Turn pilot misses into regression coverage before expansion.

### How to use this in AgentClash

Pick one workflow, freeze it as a challenge pack, record a baseline run, compare the candidate under the same tools and budget, then promote critical misses into CI release gates.