

Enterprise AI Agent Eval Checklist

A practical checklist for platform, security, and product leaders deciding whether an agent is ready for production.

AUDIENCE

Platform leads, AI product owners, security reviewers

USE IN

10 min

BEST NEXT STEP

Turn answers into a challenge pack and gate.

Workload Definition

- 1 Name the business workflow, user persona, risk level, and owner.
- 2 Convert real incidents, support tickets, or repo tasks into a versioned challenge pack.
- 3 Separate must-pass cases from exploratory cases before the first run.

Baseline And Candidate

- 1 Record the current production agent, model alias, tool policy, and harness version as the baseline.
- 2 Run candidates against the same workload, tools, time budget, and retry policy.
- 3 Compare against baseline on correctness, cost, latency, and evidence completeness.

Evidence And Review

- 1 Require replay links, scorecards, artifacts, and redacted traces for release review.
 - 2 Label evidence tier when vendor agents or hosted tools hide trajectory details.
 - 3 Keep a named reviewer for platform, security, product, and finance signoff.
-

Gate Policy

- 1 Block on correctness regression, policy violation, unsafe tool use, or hard cost and latency ceilings.
- 2 Warn on unstable exploratory metrics while the benchmark matures.
- 3 Promote every production failure into the challenge pack within one sprint.

How to use this in AgentClash

Pick one workflow, freeze it as a challenge pack, record a baseline run, compare the candidate under the same tools and budget, then promote critical misses into CI release gates.